


Beyond the Chatbot

Grounding AI in the Law

John A. Kelly


MCLE: Legal Ethics and Technology



A Process Problem, Not a Technology Problem


The electric motor took a generation to pay off, because factories kept the old layout

1890s



New tool, same layout — productivity barely moved.

1920s



New process — transformed results.

THE SAME CHOICE, TODAY



Open chatbot.
Paste question.
File what comes back.

NEW TOOL, OLD PROCESS



Ground.
Verify.
Document.

NEW PROCESS

AI hallucination is a consequence of unchanged process.
Fix the process and you fix the hallucination.

Where Courts and Practitioners Stand

6

L.A. judges in the Learned Hand pilot

ON THE BENCH

AI drafts tentative ruling



Judge edits the AI's analysis

No disclosure to the parties required.

10

states using the program

17-33%

hallucination rate in commercial legal AI (Stanford 2025)

1 in 3-6

citations wrong in an AI brief at that rate

AT COUNSEL TABLE: THE SANCTIONS TIMELINE

Mata v. Avianca, Inc.

(S.D.N.Y. 2023)
678 F.Supp.3d 443

6 fabricated opinions;
\$5,000 sanction

2023

Noland v. Land of the Free, L.P.

(2025) 114 Cal.App.5th 426
[336 Cal.Rptr.3d 897]

21 of 23 quotations
fabricated; \$10,000,
the largest in California

2025

Fivehouse v. U.S. Dept. of Defense

(E.D.N.C. 2026)
No. 2:25-cv-00041

30-year prosecutor
resigned after an AI-
drafted brief with
inaccuracies

2026

United States v. Farris

(6th Cir. Apr. 3, 2026)
2026 WL 915082

Westlaw's CoCounsel,
a commercial legal tool,
hallucinated anyway

2026

Every one used AI as a faster version of the old process: ask for cases, take what comes back, file it.

You need to know what the AI will say about your brief before the court's AI sees it.

Six Duties, New Failure Modes

Not new obligations. New ways to violate them.

COMPETENCE

Rule 1.1; State Bar Practical Guidance (2026)

Knowing where the tool fails is the competence. Ask, receive an essay, file: that is not professional judgment.

CONFIDENTIALITY

Rule 1.6; Bus. & Prof. Code § 6068(e)(1)

Every prompt is a disclosure to the vendor. A consumer chatbot with training on is no place for a client's facts.

CANDOR

Bus. & Prof. Code § 6068(d); Rule 3.3

A fabricated citation in a filed paper cannot be sat on. The fabrication was the model's; the filing is yours.

SUPERVISION

Rules 5.1, 5.3

Staff adopted AI quietly, with client data. A written office protocol is how the duty is met.

FEES

Rule 1.5; State Bar Practical Guidance (2026)

Bill the time you spent prompting and reviewing. Never bill the hours the tool saved. Put AI costs in the fee agreement.

COMMUNICATION

Rule 1.4; State Bar Practical Guidance (2026)

Tell the client when AI use materially affects scope, cost, or the method of the work. Check the client's own AI guidelines first.

THE STANDARD OF CARE IS MOVING

Electronic research became the baseline, about 2010

AI verification, the same jump, 2026

Verification is the new discipline: the output looks right by design, so the check must be mechanical.
If a tool would have caught the error and you did not use it, that is a competence problem.

The fabrication was the model's. The filing is yours.

Where Your Data Goes

Confidentiality: the webmail lesson, the court order, the fix

THE WEBMAIL LESSON

Early web mail: 40,000 unread messages left on the server. A gold mine for the provider, and reachable by subpoena to the provider. You'd never know.

Your chat history is that corpus again. The vendors want the data; they don't really want you to delete.

Digital hygiene and disposition policies apply to AI sessions like anything else in the file.

IT ALREADY HAPPENED, BY COURT ORDER

NYT v. OpenAI (S.D.N.Y. 2025):

a federal court ordered OpenAI to preserve ALL consumer output logs, including chats users had deleted, and temporary chats, under litigation hold.

The order ran May to September 2025. Users learned weeks after it began.

Discovery reaches the vendor, not just you. The vendor's records are the record.

THE FIX, SET AT THE HARNESS

ZERO DATA RETENTION

API endpoints under a ZDR agreement: inputs and outputs never logged, never retained. ZDR traffic was exempt from the preservation order; there was nothing to preserve.

Standard tiers: deleted chats purge in about 30 days, absent a legal hold.

Professional tier, training OFF, ZDR where offered.

2026 caveat: the newest frontier models keep a 30-day safety-retention window even under ZDR. Vendor policy, security-driven; flagged content held longer.

Keep a record of what you did: your work log. Don't leave the record on the vendor's cloud.

Courts and Disclosure: Three Layers, All Moving This Month

Local rules in six counties, a statute on the Governor's desk, and the Rules of Professional Conduct

LOCAL RULES: SIX COUNTIES, ONE MANDATE

Santa Clara, Gen. Rule 10 (7/1/2026):
disclose AI use in ex parte material **MANDATORY**

San Benito, Rule 10.3 (1/1/2026):
full responsibility; verify before filing

San Mateo, Rule 2.14 (1/1/2026):
follow the court's posted AI policy

Napa, Rule 3.10 (1/1/2026):
filer certifies every citation; not delegable

Sacramento, Rule 1.28 (7/1/2026):
personally read and verify every citation

Mendocino, Rules 1.54, 1.55 (7/1/2026):
rigorous human review of every filing

Alameda: no AI local rule as of July 2026.
Check department standing orders.

THE STATUTE: SB 574, ON THE GOVERNOR'S DESK

Passed 39 to 0 and 75 to 0; enrolled Aug. 31, 2026; the Governor has until Sept. 30.

New Bus. & Prof. Code § 6068.1 would:

- ✓ Bar delegating the practice of law to generative AI
- ✓ Bar entering confidential or personal identifying information into any system others can reach
- ✓ Require verifying outputs and correcting hallucinated content
- ✓ Require disclosing generative AI use to the court for all documents submitted

Code Civ. Proc. § 128.7 amended: no citation the responsible attorney has not personally verified.

THE RULES: COPRAC COMMENTS, SEPT. 11, 2026

Proposed comments to rules 1.1, 1.4, 1.6, 3.3, 5.1, and 5.3; second comment round closed; committee action Sept. 11; Supreme Court adoption required.

Rule 3.3, new comment: verify the existence and accuracy of every cited authority, AI-assisted or not.

Rule 1.4, new comment: tell the client when AI use materially affects scope, cost, or how the work is done.

The courts' own AI use: Rule of Court 10.430 and Standard 10.80, effective 9/1/2025. No disclosure to litigants required.

THE PROTOCOL: TWO LAYERS

Once per case: attach your written AI practices to the case management statement, form CM-110.

Every filing: the § 128.7 signature, and courtesy copies of every cited authority.

Tell the court your protocol before a defect does. **If SB 574 is signed, telling the court stops being optional.**

The Model Needs a Harness

Even AI's leading skeptic says the value is the scaffolding around the chatbot, not the chatbot

WHO HE IS, AND WHY TRUST HIM

Gary Marcus

- ✓ Cognitive scientist; professor emeritus of psychology and neuroscience, NYU
- ✓ Founded an AI company, Geometric Intelligence, acquired by Uber
- ✓ **Warned** LLMs would hallucinate years before ChatGPT, and was right
- ✓ **Predicted** the **Eliza effect**: we instinctively trust a talking computer like a person

“They call it the harness, but that harness is really **symbolic AI**.”

“The neural networks are pretty bad at sticking to facts. So that’s the hallucination stuff ... and the symbolic stuff, they don’t make stuff up that way.”

Gary Marcus, in conversation with Brian Greene, World Science Festival (May 2026)

STRUCTURAL BLIND SPOTS: IT PREDICTS TEXT, IT DOES NOT REASON

Not a chess engine: a 1979 Atari 2600 cartridge beat ChatGPT at chess in 2025. IBM’s Deep Blue, symbolic AI, beat the world champion. The harness marries the two.


Not a calculator: ask it long arithmetic and it estimates. Give it a calculator as a tool and it is exact.

Not a colleague: a computer that talks like a person gets trusted like a person, which is exactly why chatbot output feels more reliable than it is.

The fix is not a smarter chatbot. It is the harness around it: **grounding, verification, documentation.**

Four Types of Authority

AI hallucinates in all four, the discipline is the same



Same discipline for all four: ground the AI in sources, verify mechanically.

The Grounding Principle

A statute too new to know

**Ungrounded:
training data
only**



Training data
(cutoff
months ago)

Pattern-matches
what statutes
usually say

Misses or garbles
Veh. Code
§ 21455.9

SB 720 (Stats. 2025, ch. 782) took effect January 1, 2026: a new civil framework for camera citations. A model reasoning from memory cannot know it.

Hero statistics

17-33%

hallucination, ungrounded
commercial tools
(Magesh et al., Stanford 2025)

**Grounded:
current
statutory text
provided**



Vehicle Code,
current text,
from the repo

Distinguishes the
old § 21455.5
framework from
the new § 21455.9

Extracts every
agency obligation,
quoted with
location

From an actual matter: the traffic court case study in your materials. It ends not guilty.


Hero statistics

≈ 0%

grounded in the actual text;
most errors are retrieval
failures, not reasoning failures

Fix the inputs. The outputs follow.

The California Codes Pipeline



If the authoritative text is public, download it and ground the model. You get **this link today**.

Few-Shot Citation Prompting

Teach the model your style manual by example

WITHOUT EXAMPLES

People v. Sanchez, 63 Cal.4th 665 (2016)



Bluebook habit — year at the end, no parallel cite. Wrong manual for California briefs.

EXAMPLES PROVIDED TO THE MODEL

Sargon Enterprises, Inc. v. USC (2012) 55 Cal.4th 747 [149 Cal.Rptr.3d 614, 288 P.3d 1237]

Lazar v. Superior Court (1996) 12 Cal.4th 631 [49 Cal.Rptr.2d 377, 909 P.2d 981]

Aguilar v. Atlantic Richfield Co. (2001) 25 Cal.4th 826 [107 Cal.Rptr.2d 841, 24 P.3d 493]

California Style Manual: year in parentheses after the names; official and parallel cites bracketed.
The parallel cite is a built-in cross-check — nothing slips through.

WITH EXAMPLES

People v. Sanchez (2016) 63 Cal.4th 665 [204 Cal.Rptr.3d 102, 374 P.3d 320]



Works for any jurisdiction, any style manual.

Provide the pattern; the model follows it.

Three examples beat three paragraphs of instructions.

Show, don't tell. 3-5 examples set the format.

The 58-County Survey

One attorney, evenings and weekends, about two weeks

58

counties with their
own local rules

96

PDF files collected
and OCR'd

148

MB of rule
text

57

comparisons
vs. Sonoma

43

minutes,
full run

35:1

two-stage synthesis
(175K to 5K tokens)

10

ranked
recommendations

They differ on law and motion, filing, discovery, case management, ADR, ex parte, e-filing, tentative rulings, trial setting, and fees. No lawyer, committee, or agency had ever synthesized all 58.



58 sets of rules

ADOPTION TIERS

Widespread: 30+ counties

e-filing, tentative rulings, vertical assignment

Moderate: 14 to 22

exhibit bookmarking, disposition tracks

Innovative: 6 to 9

online reservation with forfeiture, discovery conference

Novel: 5 or fewer

no case off calendar, default fee schedule

HOW IT RAN: NO CHATBOT INVOLVED

Straight through the API: programmatic calls from the command line, all 57 comparisons fired at once. Each produced a structured report with rule citations.

The sleeper win: the model often gave the exact citation of the parallel rule in the other county, cross-jurisdiction mapping no one had.

The information was public all along. Nobody could read 148 MB.

The information existed. Nobody had time to read it all. That is exactly what the pipeline is for.

What the Model Got Wrong, and What Judgment Supplied

MODEL OUTPUT		AFTER ATTORNEY REVIEW
Sonoma lacks tentative rulings	→	Local rule 5.5 operates as auto-adoption
6 counties = widely adopted	→	6 counties = innovative (precise tier)
No preemption analysis	→	Full preemption assessment, Cal. Rules of Ct., rule 3.20

WHAT AI CONTRIBUTED	WHERE JUDGMENT WAS INDISPENSABLE
● Scope: all 58 jurisdictions ● Structure: 10 categories, consistent format ● Coverage: 78 distinct practices ● Speed: 43 minutes for 57 comparisons	● Accuracy: correcting factual errors ● Preemption: Cal. Rules of Ct., rule 3.20 analysis ● Interpretation: existing-practice context ● Terminology: field-specific precision

Comprehensive coverage is not accuracy. Let the AI handle breadth; supply the judgment; document the corrections.

The Same Discipline: One \$301 Ticket

Case study, smallest scale

THE MATTER

Red-light camera, right turn on red, Napa County. The camera photographs a man driving; the citation issues to the registered owner, who is not a man.

She checks Veh. Code § 40520: a private owner has no duty to identify the driver. She answers only “not me.”

166 days later a second citation issues, on a basis no document states. The Final Notice demands \$301 and omits the violation date entirely.

Pay, appear, or contest by written declaration (§ 40902)?

FOUR DEMONSTRATIONS, EACH USING A PROGRAM RESOURCE

1 • Codes repo

Vehicle Code grounded; every agency obligation extracted, quoted, located, including brand-new § 21455.9 (SB 720, effective 1/1/26)

2 • Local rules repo

Napa’s procedures answered from the actual rules: how the traffic clerk takes the filing

3 • Westlaw check

Every authority authenticated against the service, the Goldsmith opinion and the statute string, then loaded into the closed corpus

4 • Charging documents

The file grounded as the only source: no violation date, no issuing-agency contact, no stated basis for naming the second recipient, three different names

NOT GUILTY

Trial by written declaration decided June 30, 2026. Bail refunded in full.

THE GOVERNMENT-SIDE FLIP

Every defense ground is an agency compliance obligation. Run the same audit from the other direction, before a defendant does.

Economically irrational at ordinary drafting speed. Grounded: a few hours, and nothing went out unverified.

Case Law Is Different

STATUTES ✓

bulk download available — ground the model directly

CASE LAW ✗

paywalled, no bulk access — you cannot ground what you cannot download

THREE FAILURE MODES

Fabricated case — does not exist

✓ easy to catch: a 10-second lookup

The model invents a case name and citation (e.g., “*Smith v. Jones*, 123 F.3d 456 (2023)”). A simple database search confirms it is fake.

Real case — fabricated quotation

⚠ moderate: full-text search finds it

The model correctly cites a real case (e.g., “*Brown v. Board of Education*”), but includes a plausible-sounding quotation that does not appear in the actual opinion. Full text analysis reveals the mismatch.

Real case — mischaracterized holding

⚠ dangerous: passes the quick checks

the case is real, the quote is real, but it does not stand for the proposition cited. The model uses a dissenting opinion or dictum as the binding holding, misleading the user on the law.

The third failure mode is the one that gets through.

The Verification Ladder

Four steps, increasing rigor; each catches what the last one missed

1

Does the case exist?

Westlaw or Lexis pull; 10 seconds

every sanctioned
attorney failed HERE

2

Does the quotation exist?

full-text search of the opinion; 30 seconds

3

Good law for THIS proposition?


headnotes and Key Numbers; the step practitioners skip

4

Still good law?

KeyCite or Shepard's on the specific point

Opinion: 30 pages



HEADNOTES

[1] Negligence:
Duty of Care

[2] Evidence:
Expert Testimony

[3] Damages:
Economic Loss

[4] Procedure:
Summary Judgment

AI says the case is about economic loss. One headnote matches. Jump straight to its bracket in the text: three paragraphs, not thirty pages.

Each step catches what the previous one missed.

The Treatise and the Linter

A closed corpus and a mechanical check

The closed-corpus instruction: provide the treatise (Rutter, Witkin, the leading secondary source) and instruct: cite ONLY authorities appearing in this source. **Anything outside the corpus is, by definition, a hallucination.**

AI OUTPUT

Smith v. Jones (2019) 45 C.A.5th 123
Doe v. Roe (2021) 68 C.A.5th 456
Lee v. Park (2018) 22 C.A.5th 789
Chan v. Ito (2020) 51 C.A.5th 210

illustrative examples

THE LINTER

String comparison against
the source material. Written
FOR you, BY the harness.
You never touch code.

RESULTS, BEFORE YOU READ A WORD

- ✓ Smith v. Jones: found in source
- ✓ Doe v. Roe: found in source
- ✓ Lee v. Park: found in source
- ✗ Chan v. Ito: **FLAGGED, not in treatise**

Mechanical. Automatic. Your review time goes to the citations that passed.

Not 98% yes and 2% maybe. Yes or no.

Raw Materials and the AI Sandwich

Your own documents: the fourth category, and the biggest blind spot

WHAT IS IN YOUR FILE



It fabricates deposition quotes, invents contract clauses, misattributes statements.

THE AI SANDWICH



Two AIs on either side of a mechanical check, and they can only agree THROUGH the deterministic step. Level 1 catches fabrication; level 2 catches mischaracterization.

What a headnote does for case law, the sandwich does for any document in your file: symbolic logic, in a harness, exactly as Marcus said.

Demonstration Four of the traffic case ran this against the charging documents, and won.

Most real-world hallucination happens here, in your own file, and nobody is talking about it.

The Harness

Grounding in, verification out, documentation always

LAYER 1: GROUNDING

Codes pipeline
statutes

Rules repo
58 counties

Treatise corpus
case law

File access
raw materials

the chatbot

one small component: the interface, not the value

LAYER 2: VERIFICATION

Citation linter
string match

AI sandwich
characterization

Verification ladder
case law

Attorney review
judgment

The harness runs on YOUR computer with YOUR files. Only the model is remote.

LAYER 3: DOCUMENTATION

Session log, the flight data recorder:
automatic, JSON, on your machine

Work log, the yellow legal pad: two AI-written sentences per session; billing, handoff, work product

Claude Code · Codex · Cursor, and legal-dedicated tools are coming up fast.

A chatbot gives you a text box. A harness gives you a pipeline.

The Vendor Question

Buying advice: two structural problems

Problem 1: they trail the frontier

A generation behind: an early \$500-a-month legal product ran the prior generation while a \$20 frontier subscription was two generations ahead. Better output, a fraction of the price.

Overpromised: CoCounsel promised whole memoranda, “the rainbow of law.” It was much closer to Tesla’s so-called full self-driving.

Stale fast: the cycle runs month to month; today’s comparison is stale in a quarter. Evaluate architecture, not feature lists.

Problem 2: the black box

Reveal Codes: WordPerfect 5.1 had F5. You saw exactly what the processor was doing, and you fixed the document yourself. Word took the codes away and left you guessing at 2 a.m.

The Word thing: vendor AI products are abstracting the mechanics away and handing back an essay. Lawyers need to see the codes.

A local harness with a frontier model is Reveal Codes for AI: the narrative of every step, on your machine.

In fairness: responsible vendors exist, and for the business of law (billing, calendaring, deadlines) the products are genuinely strong today. **Don’t cancel your Westlaw or Lexis accounts.** But know what you’re buying.

The buyer's test: no symbolic-logic layer around the LLM output? You are at sea.

Seven Things This Week

Concrete, actionable, in your handout

- 1 Stress-test your next brief with AI before filing.** Your opponent will, and the court's AI may already.
- 2 Verify every AI citation** up the ladder: existence, quotation, headnotes, citator.
- 3 Provide the source material:** the code, the treatise, the transcript. Never rely on training data.
- 4 Professional tier; training off; ZDR** where offered. Your AI sessions are attorney work product.
- 5 Start a work log, one file per matter;** the AI writes it. Then drop it into your billing program: two birds.
- 6 Write the office AI protocol,** and tell the court: attach your verification practices to the CMC statement.
- 7 File courtesy copies of every cited authority.** The court already holds every case you cite. Costs nothing.

Start here. The checklist is in your handout.

What You Take Home

Three electronic resources, links live now

All 29 California Codes

machine-readable, nightly sync from the Legislature's bulk data

github.com/johnakelly-yahoo-com/california-codes

Usage: Point the model at the Probate Code file and instruct: answer only from this text.

All 58 County Local Rules

structured for AI retrieval — full rules plus civil-only extracts

github.com/johnakelly-yahoo-com/california-local-rules

Usage: Load your county's file before drafting anything the local rules govern.

Tested Prompts

case law verification, citation checking, document-grounded research — designed for Westlaw and Lexis

in your handout materials

Usage: Run the four-step verification prompt on every AI-drafted brief.

Links live now. Implement this week.

“Beyond the Chatbot”

Grounding AI in the Law

 Ground it in the sources.

 Verify it mechanically.

 Document what you did.

The method of finding the law shapes what law you find.

See generally Robert C. Berring, *Legal Research and Legal Concepts: Where Form Molds Substance* (1987) 75 Cal. L.Rev. 15

The chatbot offered this line as a direct quotation. I checked the article; the sentence appears nowhere in it. Hence the *see generally*. The signal is doing its job.